AI in Ramadan

Session 2 : Diving into abyss

By Loukmane MAADA

Recap

AI approaches

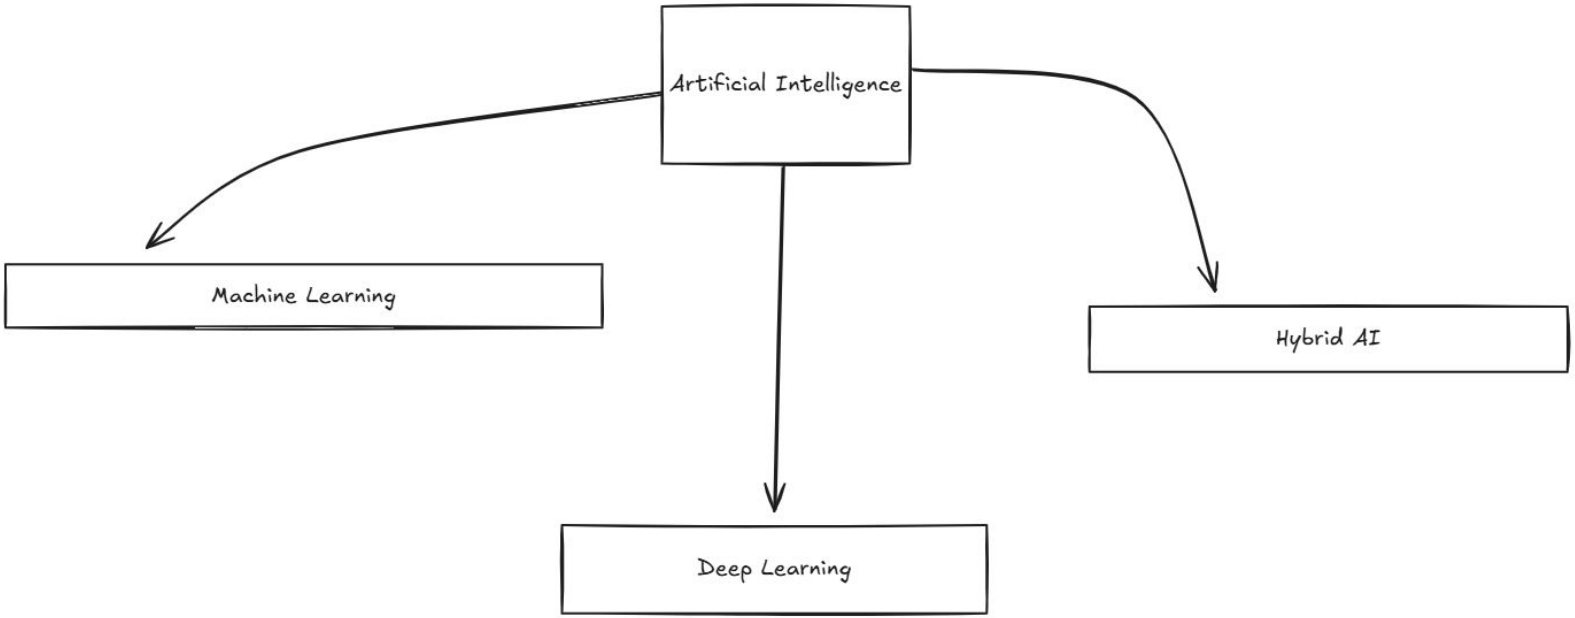

AI workflow

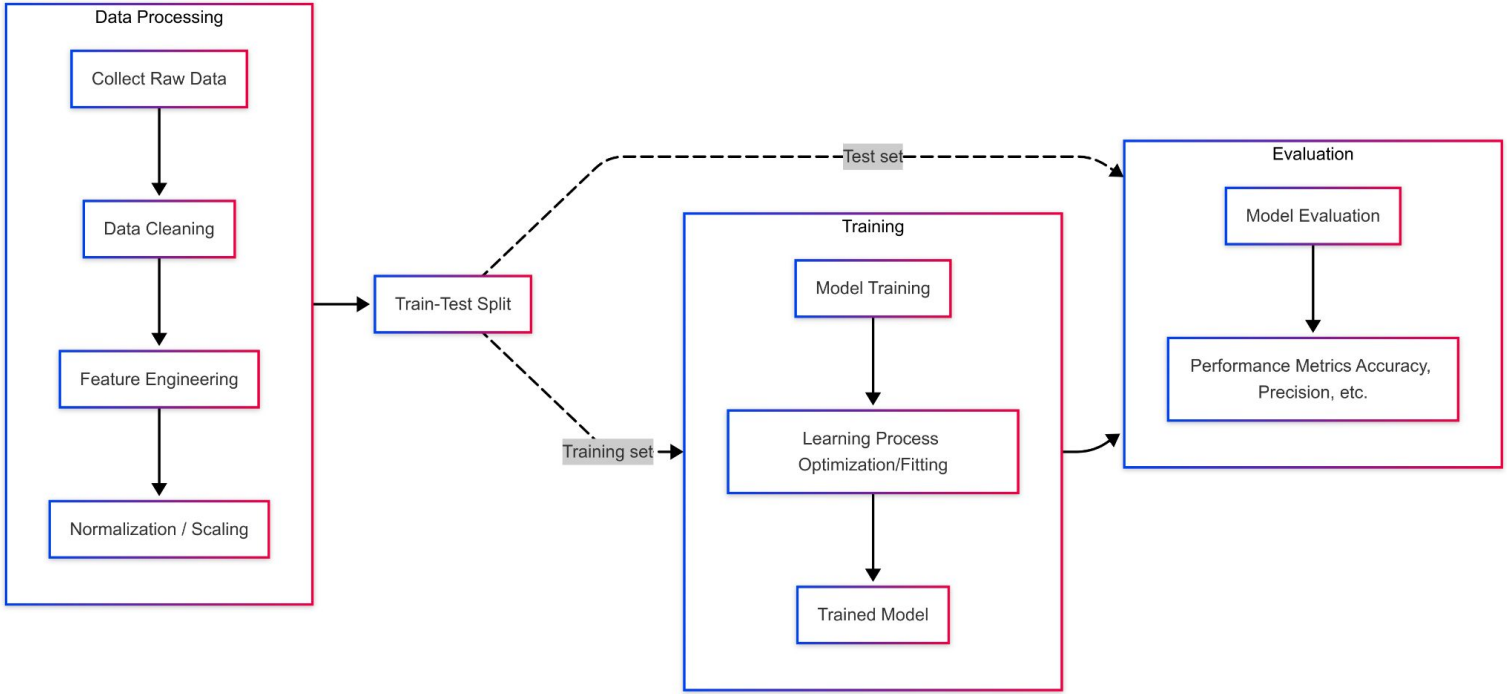
Data Processing
Collect Raw Data
Data Cleaning
Feature Engineering
Normalization / Scaling
Train-Test Split
Test set
Training set
Training
Model Training
Learning Process
Optimization/Fitting
Trained Model
Evaluation
Model Evaluation
Performance Metrics Accuracy,
Precision, etc.

Introduction

Introduction

A subset of AI that uses neural networks with multiple layers (hence "deep") to automatically learn representations from data. Deep learning (DL) eliminates much of the manual feature engineering by learning layered feature hierarchies directly from raw data.

Why DL ?

Deep learning approaches have achieved state-of-the-art results in many domains. They can automatically extract intricate patterns and relationships in data through deep neural network architectures.

Examples of DL in real life

- Computer Vision: Image classification AlexNet (e.g. identifying objects in photos) and facial recognition.
- Natural Language Processing: Machine translation (e.g. Google Translate), sentiment analysis of text, chatbots and virtual assistants.
- Robotics and Games: Autonomous driving systems, drones, and AI agents mastering complex games (e.g. AlphaGo, OpenAI 5).

The perceptron

Neuron

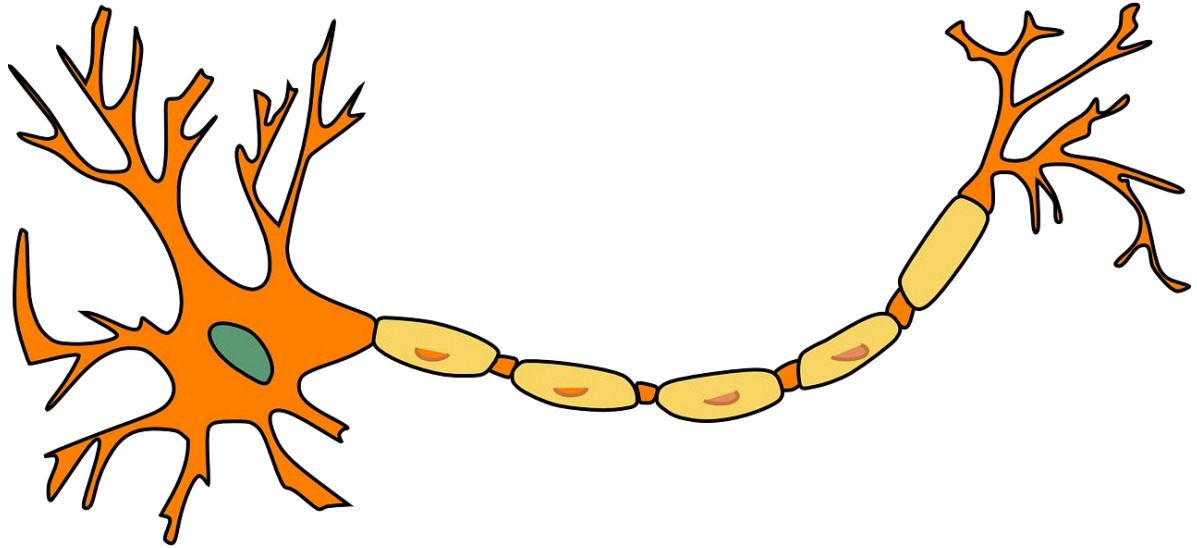

Perceptron

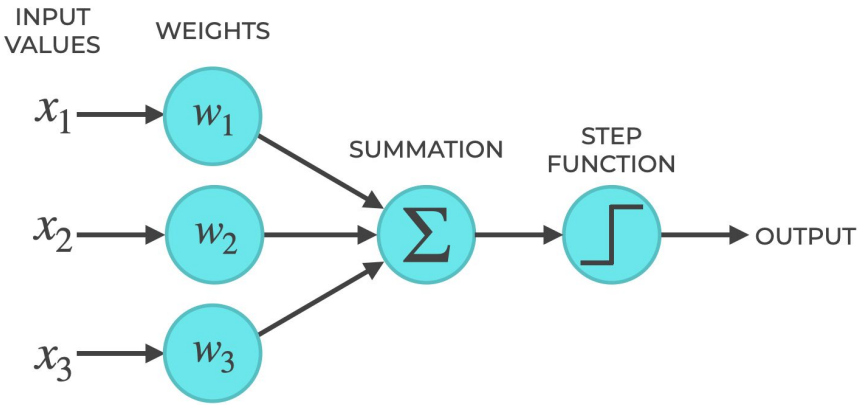
INPUT VALUES
WEIGHTS
x_1
x_2
x_3
w_1
w_2
w_3
SUMMATION
Σ
STEP FUNCTION
OUTPUT

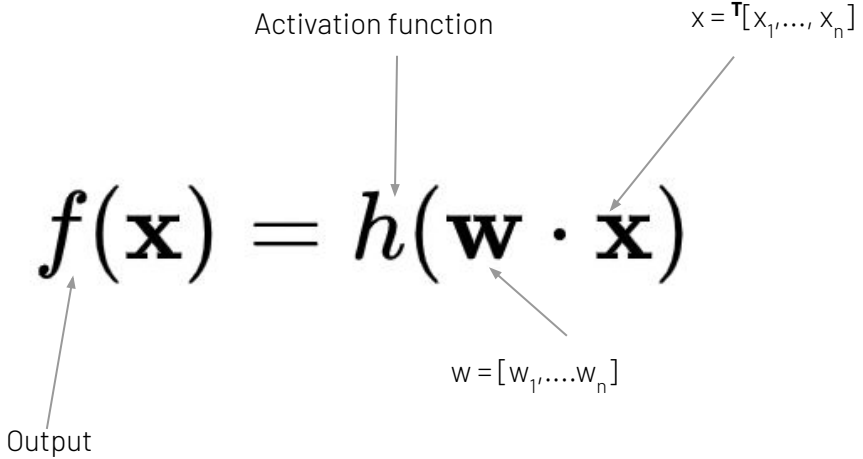
Activation function
$x = {}^{T}[x_1, \ldots, x_n]$
$f(\mathbf{x}) = h(\mathbf{w} \cdot \mathbf{x})$
$w = [w_1, \ldots . w_n]$
Output

Multilayer Perceptron

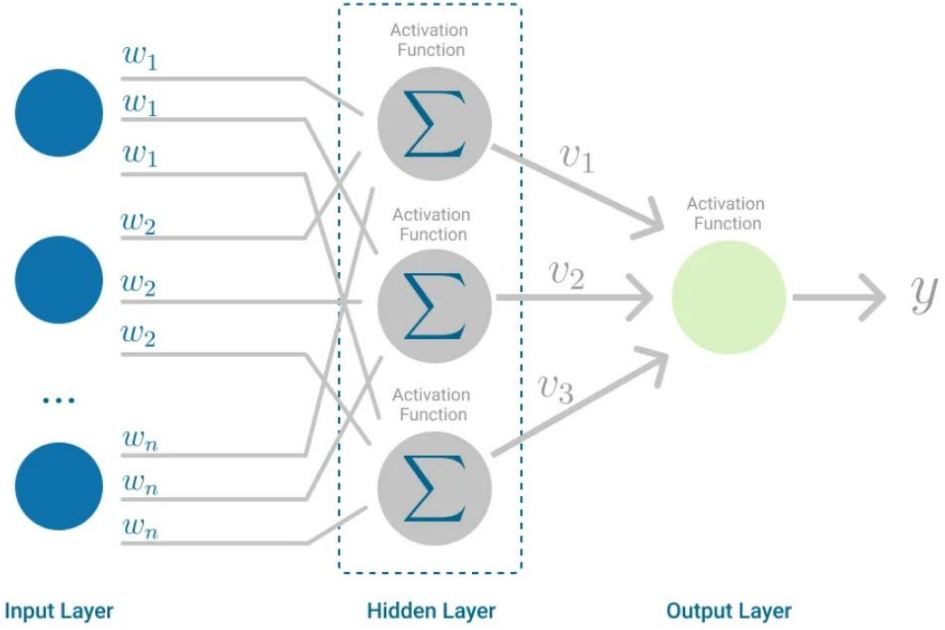

Training

Loss function

A loss (or cost) function measures the error between the network's predictions and the true targets. It's a single number that the training process tries to minimize.

Mean Squared Error (MSE)

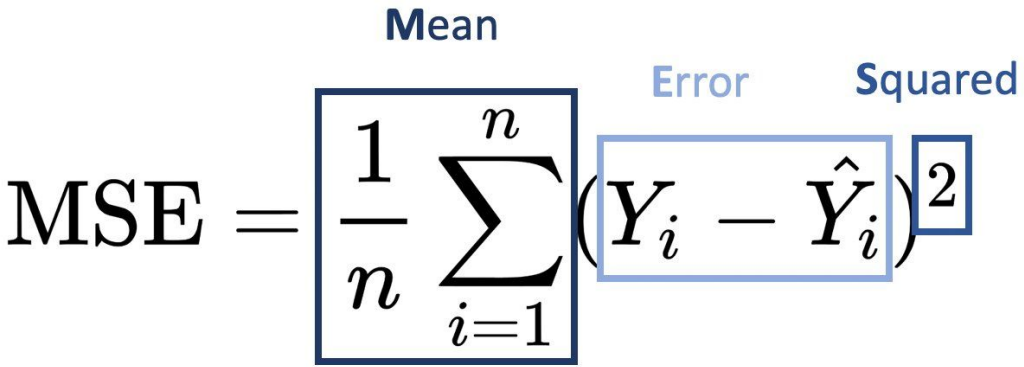

Cross-entropy loss

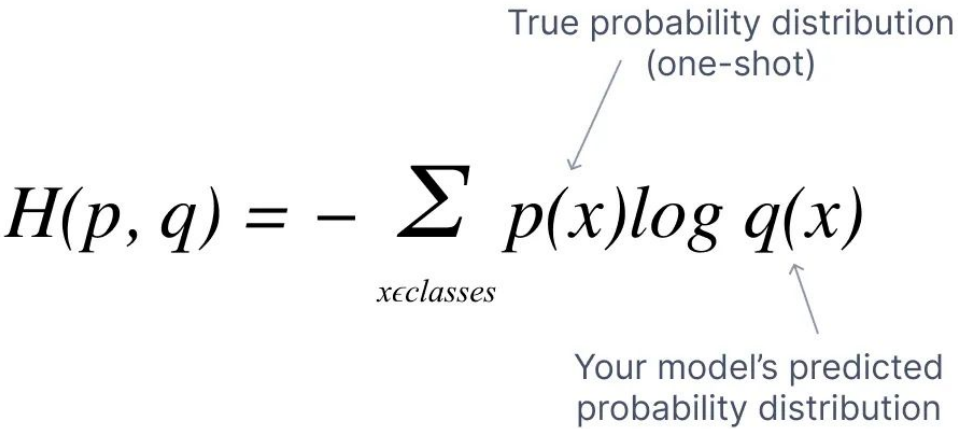

Gradient descent (backpropagation)

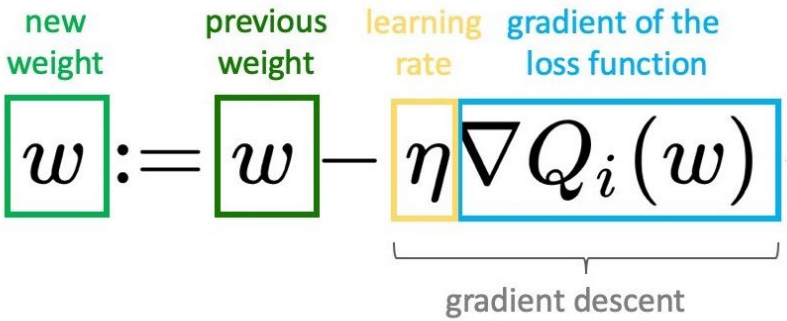

Advanced Models

CNN

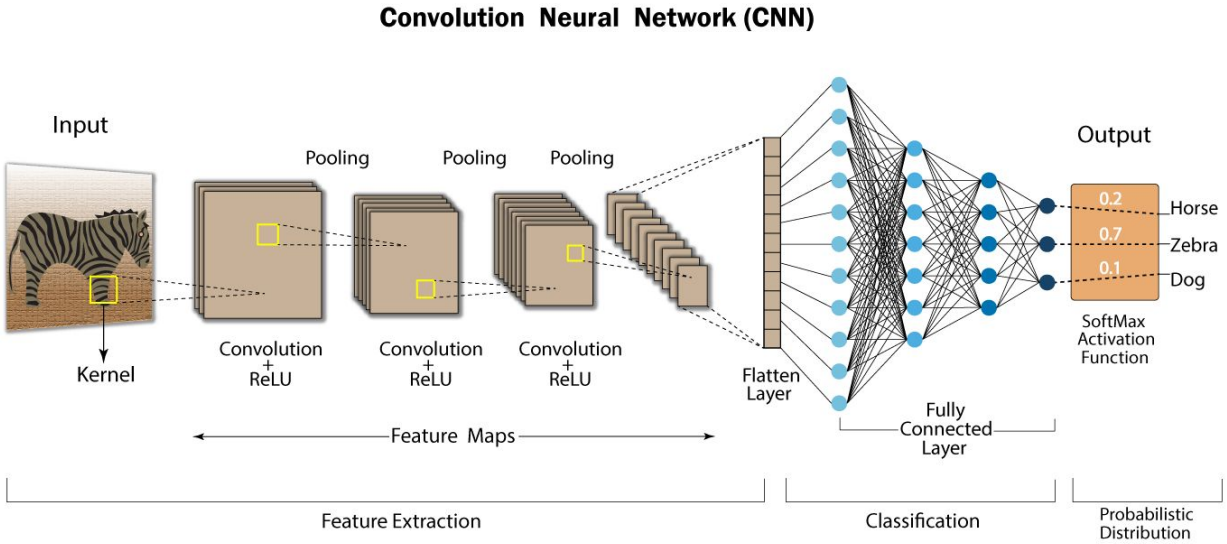
Convolution Neural Network (CNN)
Input
Kernel
Pooling
Pooling
Pooling
Convolution + ReLU
Convolution + ReLU
Convolution + ReLU
Feature Maps
Flatten Layer
Fully Connected Layer
Output
0.2
0.7
0.1
Horse
Zebra
Dog
SoftMax Activation Function
Feature Extraction
Classification
Probabilistic Distribution

Convolution

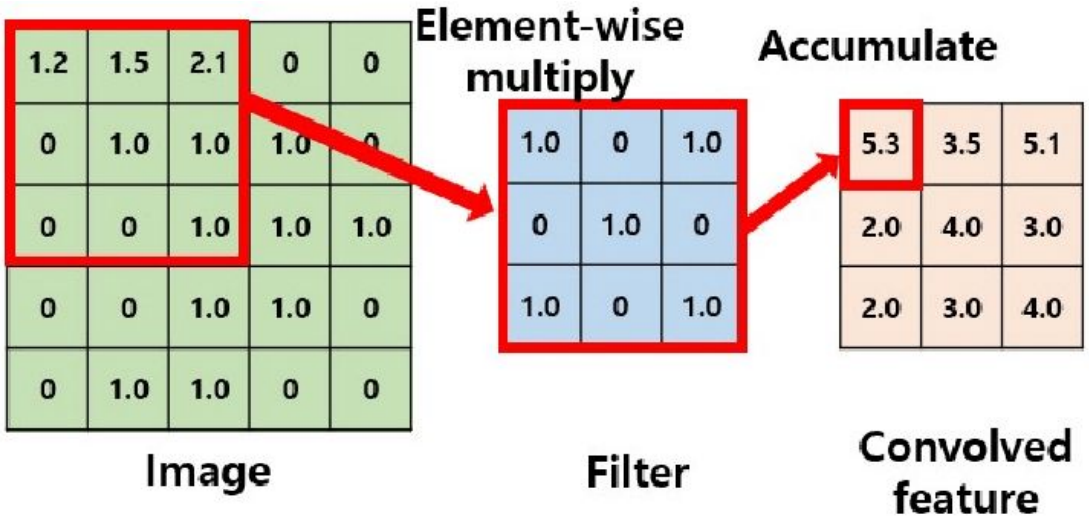

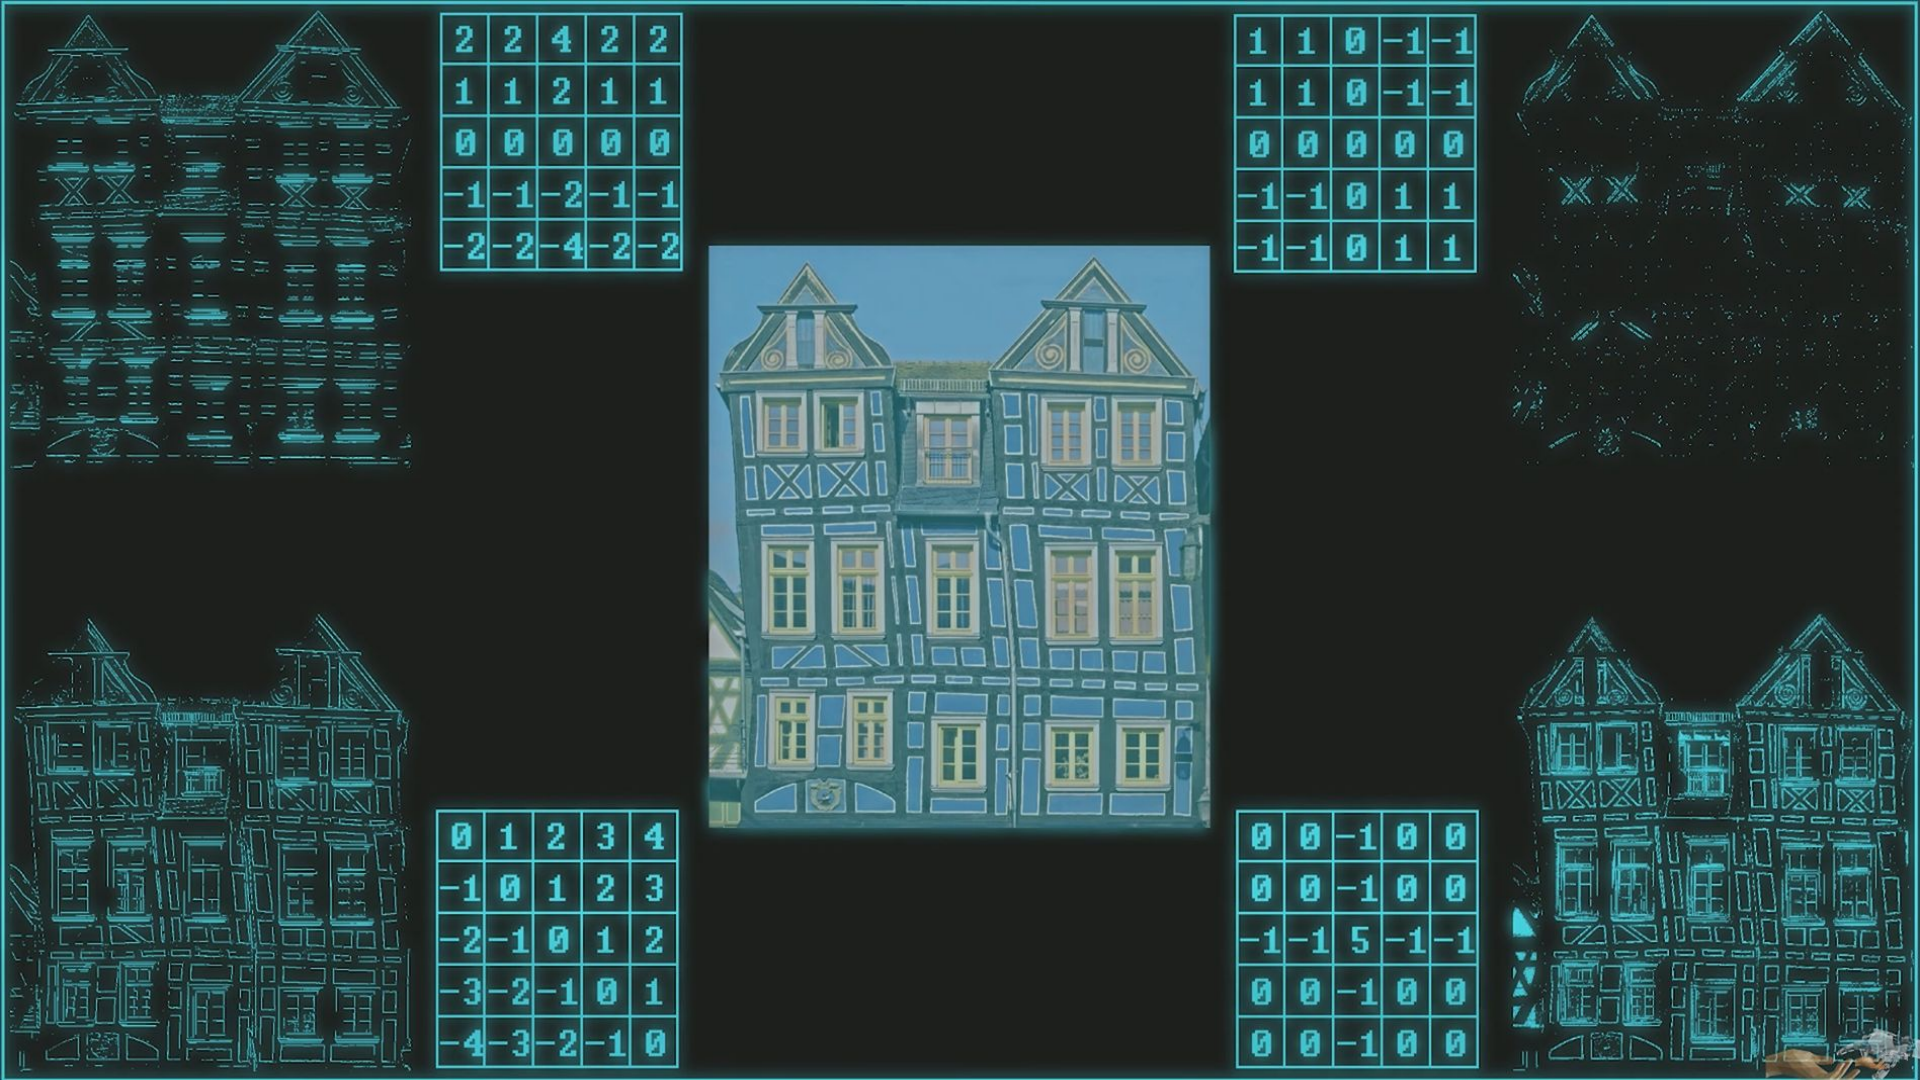

2	2	4	2	2
1	1	2	1	1
0	0	0	0	0
-1	-1	-2	-1	-1
-2	-2	-4	-2	-2

1	1	0	-1	-1
1	1	0	-1	-1
0	0	0	0	0
-1	-1	0	1	1
-1	-1	0	1	1

0	1	2	3	4
-1	0	1	2	3
-2	-1	0	1	2
-3	-2	-1	0	1
-4	-3	-2	-1	0

0	0	-1	0	0
0	0	-1	0	0
-1	-1	5	-1	-1
0	0	-1	0	0
0	0	-1	0	0

pooling

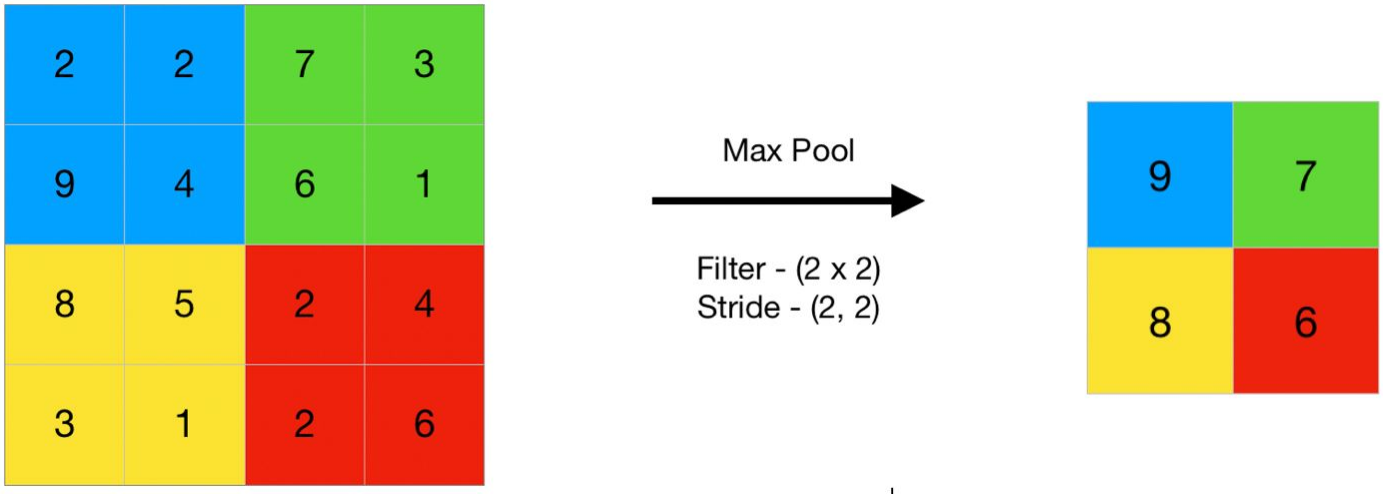

RNN

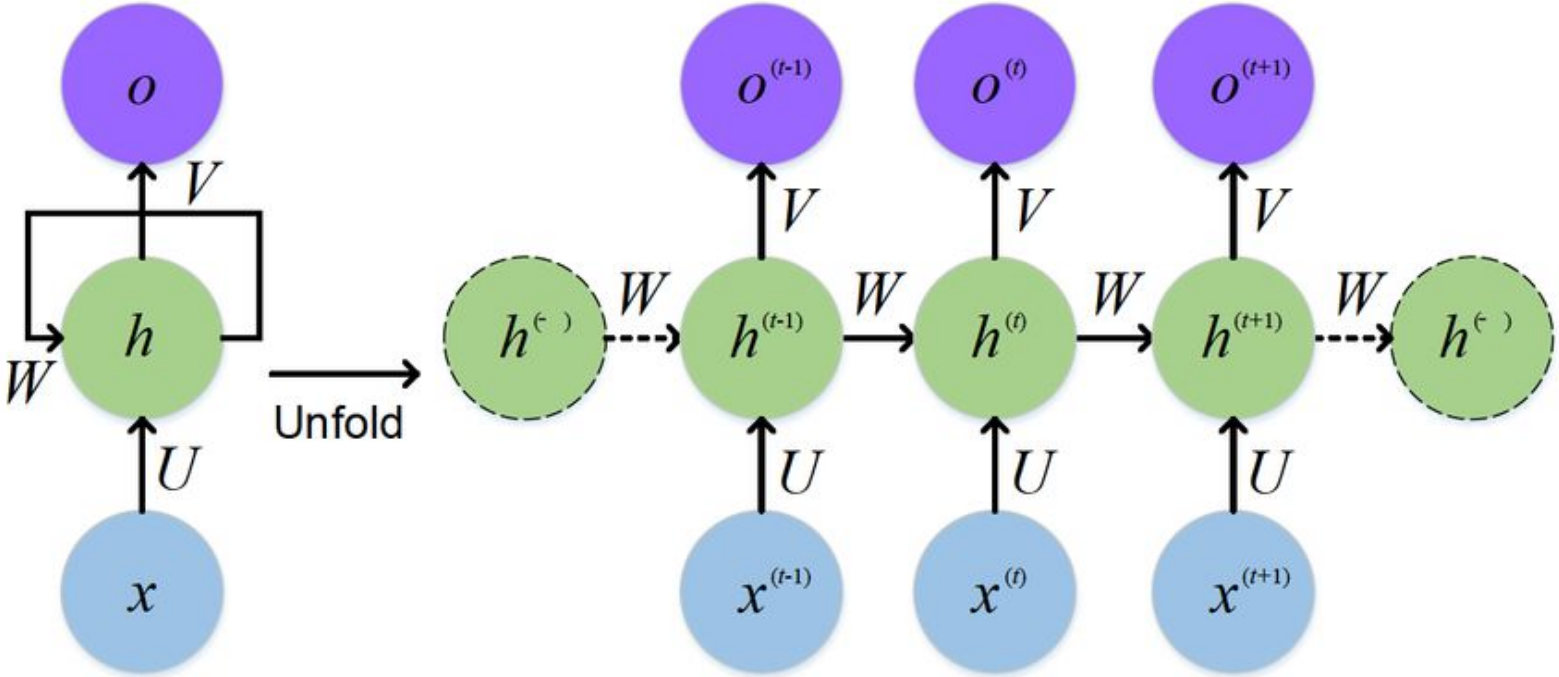

RNN heuristic

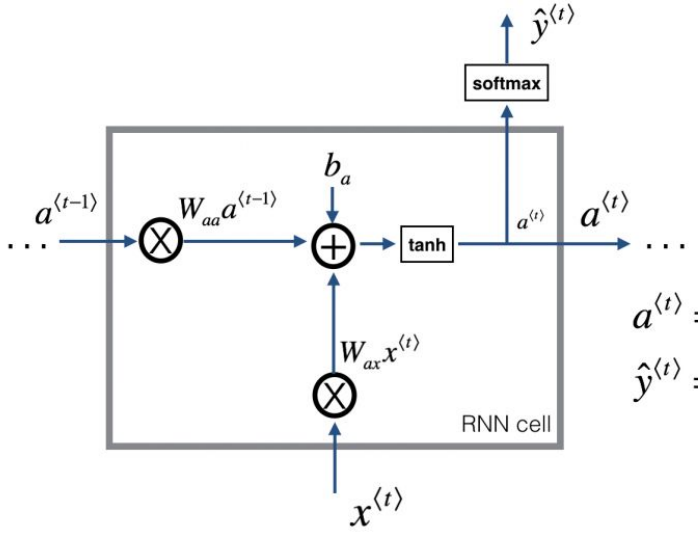

$$a^{\langle t \rangle} = \tanh(W_{ax}x^{\langle t \rangle} + W_{aa}a^{\langle t-1 \rangle} + b_a)$$

$$\hat{y}^{\langle t \rangle} = soft\max(W_{ya}a^{\langle t \rangle} + b_y)$$

LSTM Architecture

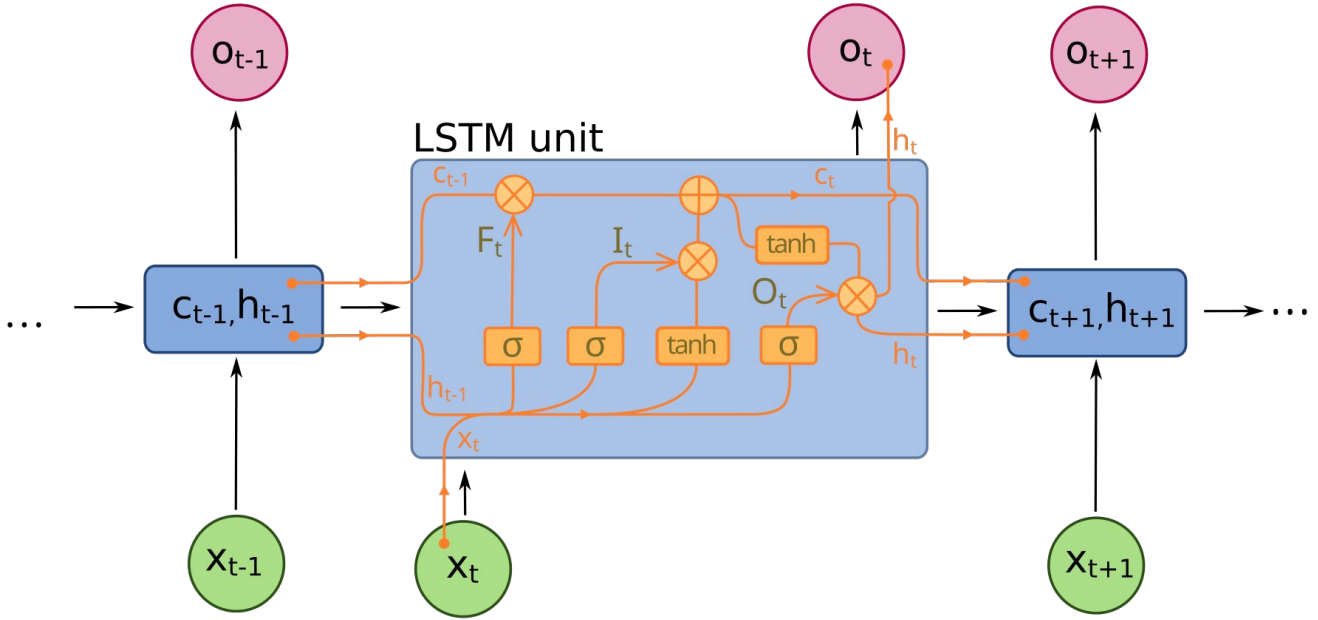

GRU Architecture

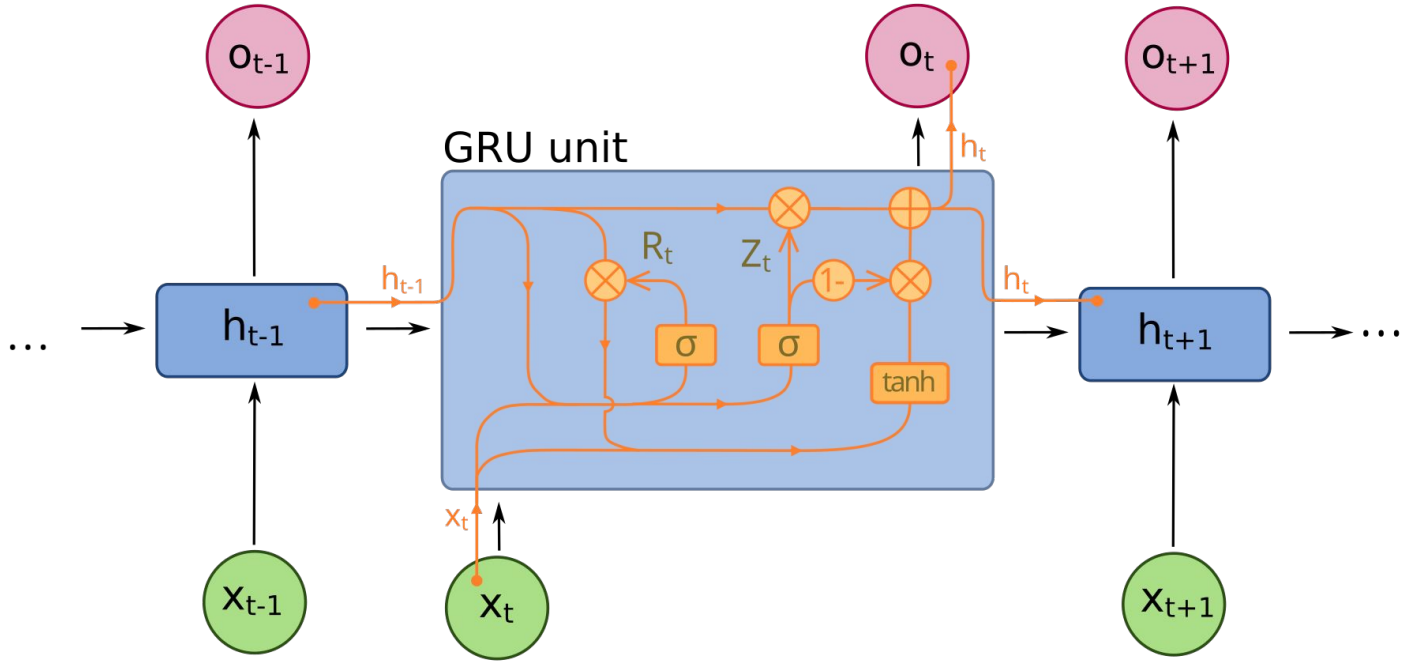

Questions?

See you

next time

つづく